NYU Center for Data Science: DS-GA 1003
Machine Learning and Computational Statistics (Spring 2019)

Brett Bernstein

March 13, 2019

Instructions: Following most lab and lecture sections, we will be providing concept checks for review. Each concept check will:

- List the lab/lecture learning objectives. You will be responsible for mastering these objectives, and demonstrating mastery through homework assignments, exams (midterm and final), and on the final course project.

- Include concept check questions. These questions are intended to reinforce the lab/lectures, and help you master the learning objectives.

You are strongly encourage to complete all concept check questions, and to discuss these (and related) problems on Piazza and at office hours. However, problems marked with a ($\star$) are considered optional.

Conditional Probability Models: Concept Check

Conditional Probability Models

MLE Learning Objectives

- Define the likelihood of an estimate of a probability distribution for some data $\mathcal{D}$.

- Define a parameteric model, and some common parameteric families.

- Define the MLE for some parameter θ of a probability model.

- Be able to find the MLE using first order conditions on the log-likelihood.

Conditional Probability Models

- Describe the basic structure of a linear probabilistic model, in terms of (i) a parameter θ of the probablistic model, (ii) a linear score function, (iii) a transfer function (kin to a "response function" or "inverse link" function, though we've relaxed requirements on the parameter theta).

- Explain how we can use MLE to choose w, the weight vector in our linear function (in (ii) above).

- Give common transfer functions for (i) bernoulli, (ii) poisson, (iii) gaussian, and (iv) categorical distributions. Explain why these common transfer functions make sense (in terms of their codomains).

- Explain the equivalence of EMR and MLE for negative log-likelihood loss.

MLE/Conditional Probability Model Concept Check Question

1. In each of the following, assume $X_1, \ldots, X_n$ are an i.i.d. sample from the given distribution.

   (a) Compute the MLE for p assuming each $X_i \sim \text{Geom}(p)$ with PMF $f_X(k) = (1 - p)^{k-1} p$ for $k \in \mathbb{Z}_{\geq 1}$.

   (b) Compute the MLE for λ assuming each $X_i \sim \text{Exp}(\lambda)$ with PDF $f_X(x) = \lambda e^{-\lambda x}$.

   Solution.

   (a) The likelihood L is given by

   $$L(p; x_1, \ldots, x_n) = \prod_{i=1}^{n} (1 - p)^{x_i - 1} p$$

   giving a log-likelihood

   $$\log L(p; x_1, \ldots, x_n) = n \log p + \left(\sum_{i=1}^{n} x_i - 1 \right) \log(1 - p).$$

   Differentiating gives

   $$\frac{d}{dp} \log L(p; x_1, \ldots, x_n) = \frac{n}{p} - \frac{\sum_{i=1}^{n} x_i - 1}{1 - p}.$$

   Solving for a critical point we get

   $$\frac{d}{dp} \log L(p; x_1, \ldots, x_n) = 0 \iff \frac{1}{n} \sum_{i=1}^{n} x_i = \frac{1}{p} \iff p = \frac{n}{\sum_{i=1}^{n} x_i}.$$

   By the first or second derivative tests, this is the maximum. Thus the answer is

   $$\hat{p}_{\text{MLE}} = \frac{n}{\sum_{i=1}^{n} x_i}.$$

(b) The likelihood L is given by

$$L(\lambda; x_1, \ldots, x_n) = \prod_{i=1}^{n} \lambda e^{-\lambda x_i}$$

giving a log-likelihood

$$\log L(\lambda; x_1, \ldots, x_n) = n \log \lambda - \lambda \sum_{i=1}^{n} x_i.$$

Differentiating gives

$$\frac{d}{dp} \log L(p; x_1, \ldots, x_n) = \frac{n}{\lambda} - \sum_{i=1}^{n} x_i.$$

Solving for a critical point we get

$$\frac{d}{dp} \log L(p; x_1, \ldots, x_n) = 0 \iff \lambda = \frac{1}{n} \sum_{i=1}^{n} x_i.$$

By the first or second derivative tests, this is a maximum. Thus the answer is

$$\hat{\lambda}_{\text{MLE}} = \frac{n}{\sum_{i=1}^{n} x_i}.$$

2. We want to fit a regression model where $Y|X = x \sim \text{Unif}([0, e^{w^T x}])$ for some $w \in \mathbb{R}^d$. Given i.i.d. data points $(X_1, Y_1), \ldots, (X_n, Y_n) \in \mathbb{R}^d \times \mathbb{R}$, give a convex optimization problem that finds the MLE for w.

Solution. The likelihood L is given by

$$L(w; x_1, y_1, \ldots, x_n, y_n) = \prod_{i=1}^{n} \frac{\mathbf{1}(y_i \le e^{w^T x_i})}{e^{w^T x_i}}.$$

Taking logs we get

$$-\sum_{i=1}^{n} w^T x_i = -w^T \left(\sum_{i=1}^{n} x_i \right)$$

if $y_i \le \exp(w^T x_i)$ for all i, or $-\infty$ otherwise. Thus we obtain the linear program

$$\text{minimize} \quad w^T \left(\sum_{i=1}^{n} x_i \right)$$

$$\text{subject to} \quad \log(y_i) \le w^T x_i \quad \text{for } i = 1, \ldots, n.$$

3. Explain why softmax is related to computing the maximum of a list of values.

   Solution. Let $x_1, \ldots, x_n \in \mathbb{R}$. Let $\mathrm{ArgMax}(x_1, \ldots, x_n)$ denote a 1-hot encoding of the argmax function:

   $$\mathrm{ArgMax}(x_1, \ldots, x_n) = \left(\mathbf{1}(\arg \max_i x_i = 1), \ldots, \mathbf{1}(\arg \max_i x_i = n) \right).$$

   Recall that softmax has the following definition:

   $$\mathrm{softmax}_\lambda(x_1, \ldots, x_n) = \frac{1}{\sum_{i=1}^n e^{\lambda x_i}} \left(e^{\lambda x_1}, \ldots, e^{\lambda x_n} \right),$$

   where $\lambda > 0$ is a fixed parameter. We claim that softmax is a differentiable approximation to ArgMax. Consider what happens when we let $x_j \to \infty$ while keeping the other values fixed. Then

   $$\frac{e^{\lambda x_j}}{\sum_{i=1}^n e^{\lambda x_i}} \to 1$$

   and

   $$\frac{e^{\lambda x_k}}{\sum_{i=1}^n e^{\lambda x_i}} \to 0$$

   for all $k \neq j$. For example, suppose $x_1 = 1$, $x_2 = -3$, $x_3 = 5$. Then

   $$\mathrm{softmax}_1(1, -3, 5) = (0.0180, 0.0003, 0.9817)$$

   while

   $$\mathrm{ArgMax}(1, -3, 5) = (0, 0, 1).$$

4. Suppose x has a Poisson distribution with unknown mean θ:

   $$p(x|\theta) = \frac{\theta^x}{x!} \exp(-\theta), \qquad x = 0, 1, \cdots$$

   Let the prior for θ be a gamma distribution:

   $$p(\theta|\alpha, \beta) = \frac{\beta^\alpha \theta^{\alpha-1}}{\Gamma(\alpha)} \exp(-\beta\theta), \qquad \theta > 0$$

   where Γ is the gamma function. Show that, given an observation x, the posterior $p(\theta|x, \alpha, \beta)$ is a gamma distribution with updated parameters $(\alpha', \beta') = (\alpha + x, \beta + 1)$. What does this tell you about the Poisson and gamma distributions?

   Solution. From Bayes' theorem[1], we have:

   $$\begin{aligned}
   p(\theta|x) &\propto p(x|\theta)p(\theta) \\
   &\propto \left(\theta^x \exp(-\theta) \right) \left(\theta^{\alpha-1} \exp(-\beta\theta) \right) \\
   &= \theta^{x+\alpha-1} \exp\left(-(\beta + 1)\theta) \right) \\
   &\propto \mathcal{G}(\alpha + x, \beta + 1)
   \end{aligned}$$

---

[1]Actually from Roman Garnett, from whom this problem was taken.

This shows that the gamma is the conjugate prior to the Poisson. Also, note here we exploit a common trick: we manipulate the numerator, ignoring constants independent of θ. If we can recognize the functional form as belonging to a distribution family we know, we can simply identify the parameters and trust that the distribution normalizes!